



# 1. Setup Minikube in your local machine.

Install docker desktop in your local machine.



Install oracle virtualbox in your local machine.



Search on browser to install minikube in windows.

**Copy this command this command and give this in powershell.**


```

New-Item -Path 'c:\' -Name 'minikube' -ItemType Directory -Force
$ProgressPreference = 'SilentlyContinue'; Invoke-WebRequest -OutFile 'c:\minikube\minikube.exe' -Uri 'https://github.com/kubernetes/minikube/releases/latest/minikube.exe'

```

- New-Item -Path 'c:\' -Name 'minikube' -ItemType Directory -Force\$ProgressPreference = 'SilentlyContinue'; Invoke-WebRequest -OutFile 'c:\minikube\minikube.exe' -Uri 'https://github.com/kubernetes/minikube/releases/latest/minikube.exe'**

# est/download/minikube-windows-amd64.exe' - UseBasicParsing



```
Windows PowerShell
Copyright (C) Microsoft Corporation. All rights reserved.

Install the latest PowerShell for new features and improvements! https://aka.ms/PSWindows


PS C:\Users\Ashish> New-Item -Path 'c:\' -Name 'minikube' -ItemType Directory -Force

    Directory: C:\


Mode                LastWriteTime     Length Name
----                <-----          ----- 
d-----        16-12-2025     18:21      minikube

PS C:\Users\Ashish> $ProgressPreference = 'SilentlyContinue'; Invoke-WebRequest -OutFile 'c:\minikube\minikube.exe' -Uri 'https://github.com/kubernetes/minikube/releases/latest/download/minikube-windows-amd64.exe' -UseBasicParsing
```


Go to environmental variables and copy the path of minikube and paste in the path



Click on edit.



Paste the copied path.



Go to visual studio and execute this command.


- **minikube start**

```
PS C:\Users\Ashish\Desktop\Terraform basics> minikube start
🌟 minikube v1.37.0 on Microsoft Windows 11 Home Single Language 10.0.26200.7462 Build 26200.7462
🌟 Automatically selected the docker driver. Other choices: virtualbox, ssh
⭐ Using Docker Desktop driver with root privileges
👉 Starting "minikube" primary control-plane node in "minikube" cluster
_PULLING_ Pulling base image v0.0.48 ...
⬇️ DOWNLOADING_ Downloading Kubernetes v1.34.0 preload ...
  > preloaded-images-k8s-v18-v1...: 337.07 MiB / 337.07 MiB 100.00% 9.76 Mi
  > gcr.io/k8s-minikube/kichbase...: 488.52 MiB / 488.52 MiB 100.00% 10.52 M
🔥 Creating docker container (CPUs=2, Memory=3900MB) ...
⚠️ Failing to connect to https://registry.k8s.io/ from inside the minikube container
💡 To pull new external images, you may need to configure a proxy: https://minikube.sigs.k8s.io/docs/reference/networking/proxy/
🌐 Preparing Kubernetes v1.34.0 on Docker 28.4.0 ...
🔗 Configuring bridge CNI (Container Networking Interface) ...
🌐 Verifying Kubernetes components...
  - Using image gcr.io/k8s-minikube/storage-provisioner:v5
🌟 Enabled addons: storage-provisioner, default-storageclass

❗ C:\Program Files\Docker\Docker\resources\bin\kubectl.exe is version 1.32.2, which may have incompatibilities with Kubernetes 1.34.0.
  - Want kubectl v1.34.0? Try 'minikube kubectl -- get pods -A'
🔥 Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
PS C:\Users\Ashish\Desktop\Terraform basics> minikube ip
192.168.49.2
PS C:\Users\Ashish\Desktop\Terraform basics>
```

## 2. Setup k8s master and two worker nodes on ubuntu.

Launch 3 instances with c7 large and memory of 20gb.



The screenshot shows the AWS EC2 Instances page. On the left sidebar, under the 'Instances' section, 'Instances' is selected. The main area displays a table of instances with the following details:

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4
K8s-master	i-05fd47928fc1b0b5f	Running	c7i-flex.large	Initializing	<a href="#">View alarms +</a>	us-east-1a	ec2-3-81-57
worker-01	i-0a25fe43cc6ad720b	Running	c7i-flex.large	Initializing	<a href="#">View alarms +</a>	us-east-1a	ec2-3-91-22
worker-02	i-0fbab4f4013855ca	Running	c7i-flex.large	Initializing	<a href="#">View alarms +</a>	us-east-1a	ec2-18-234

Below the table, the details for the selected instance, 'K8s-master', are shown. The 'Details' tab is active, displaying sections for Status and alarms, Monitoring, Security, Networking, Storage, and Tags.


Login to master instance.

- **hostname master**
- **echo \$ hostname**
- **vi /etc/hostname**

```
[root@ip-172-31-24-182 ~]# hostname master
[root@ip-172-31-24-182 ~]# echo $hostname

[root@ip-172-31-24-182 ~]# vi /etc/hostname
[root@ip-172-31-24-182 ~]# |
```

Here we are replacing our ipaddress as master for our identification.



Logout and login again you will get like this.

```
[root@master ~]# |
```

Now do for same like this for worker nodes for understanding purpose.

```
[root@ip-172-31-27-50 ~]# hostname worker-01
[root@ip-172-31-27-50 ~]# echo $hostname

[root@ip-172-31-27-50 ~]# vi /etc/hostname
[root@ip-172-31-27-50 ~]#
```

worker-01

```
[root@worker-01 ~]#
```

```
[root@ip-172-31-72-26 ~]# hostname worker-02
[root@ip-172-31-72-26 ~]# echo $hostname

[root@ip-172-31-72-26 ~]# vi /etc/hostname
[root@ip-172-31-72-26 ~]#
```

worker-02

```
[root@worker-02 ~]#
```

Install docker in all 3 servers.

- **yum install docker -y**
- **systemctl start docker**
- **systemctl enable docker**

```

Running scriptlet: docker-25.0.13-1.amzn2023.0.2.x86_64
created symlink /etc/systemd/system/sockets.target.wants/docker.socket → /usr/lib/systemd/system/docker.socket.

Running scriptlet: container-selinux-4:2.242.0-1.amzn2023.noarch
Running scriptlet: docker-25.0.13-1.amzn2023.0.2.x86_64
Verifying   : container-selinux-4:2.242.0-1.amzn2023.noarch
Verifying   : containerd-2.1.5-1.amzn2023.0.1.x86_64
Verifying   : docker-25.0.13-1.amzn2023.0.2.x86_64
Verifying   : iptables-libs-1.8.8-3.amzn2023.0.2.x86_64
Verifying   : iptables-nft-1.8.8-3.amzn2023.0.2.x86_64
Verifying   : libcgroup-3.0-1.amzn2023.0.1.x86_64
Verifying   : libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64
Verifying   : libnftnlink-1.0.1-19.amzn2023.0.2.x86_64
Verifying   : libnftnl-1.2.2-2.amzn2023.0.2.x86_64
Verifying   : pigz-2.5-1.amzn2023.0.3.x86_64
Verifying   : runc-1.3.3-2.amzn2023.0.1.x86_64

Installed:
  container-selinux-4:2.242.0-1.amzn2023.noarch      containerd-2.1.5-1.amzn2023.0.1.x86_64      docker-25.0.13-1.am
  iptables-libs-1.8.8-3.amzn2023.0.2.x86_64        iptables-nft-1.8.8-3.amzn2023.0.2.x86_64      libcgroup-3.0-1.amz
  libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64    libnftnlink-1.0.1-19.amzn2023.0.2.x86_64      libnftnl-1.2.2-2.am
  pigz-2.5-1.amzn2023.0.3.x86_64                  runc-1.3.3-2.amzn2023.0.1.x86_64

Complete!
[root@master ~]# systemctl start docker
[root@master ~]# systemctl enable docker
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
[root@master ~]#

```

Search in browser as kubeadm and install this command in 3 servers.

**cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo**

**[kubernetes]**

**name=Kubernetes**

**baseurl=https://pkgs.k8s.io/core:/stable:/v1.34/rpm/**

**enabled=1**

**gpgcheck=1**

**gpgkey=https://pkgs.k8s.io/core:/stable:/v1.34/rpm/repo/ata/repo/repodata.xml.key**

**exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni**

**EOF**

```
[root@master ~]# cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.34/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.34/rpm/repo/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.34/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.34/rpm/repo/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
[root@master ~]#
```

```
[root@worker-01 ~]# cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.34/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.34/rpm/repo/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.34/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.34/rpm/repo/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
[root@worker-01 ~]#
```

```
[root@worker-02 ~]# cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.34/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.34/rpm/repo/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.34/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.34/rpm/repo/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
[root@worker-02 ~]#
```

Paste this command in 3 servers,

- **sudo setenforce 0**

- **sudo sed -i**

```
's/^SELINUX=enforcing$/SELINUX=permissive/'  
/etc/selinux/config
```

```
[root@master ~]# sudo setenforce 0  
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config  
[root@master ~]# |
```

```
[root@worker-01 ~]# sudo setenforce 0  
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config  
[root@worker-01 ~]# |
```

```
[root@worker-02 ~]# sudo setenforce 0  
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config  
[root@worker-02 ~]# |
```

Install this in 3 servers.

- **sudo yum install -y kubelet kubeadm kubectl -- disableexcludes=Kubernetes**

```
[root@master ~]# sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes  
Kubernetes  
Dependencies resolved.  
=====  
 Package           Architecture      Version  
=====  
Installing:  
 kubeadm          x86_64          1.34.3-150500.1.1  
 kubectl          x86_64          1.34.3-150500.1.1  
 kubelet          x86_64          1.34.3-150500.1.1  
Installing dependencies:  
 cri-tools        x86_64          1.34.0-150500.1.1  
 kubernetes-cni   x86_64          1.7.1-150500.1.1  
  
Transaction Summary  
=====  
Install 5 Packages  
  
Total download size: 52 M  
Installed size: 282 M  
Downloading Packages:  
(1/5): cri-tools-1.34.0-150500.1.1.x86_64.rpm  
(2/5): kubectl-1.34.3-150500.1.1.x86_64.rpm  
(3/5): kubeadm-1.34.3-150500.1.1.x86_64.rpm  
(4/5): kubernetes-cni-1.7.1-150500.1.1.x86_64.rpm  
(5/5): kubelet-1.34.3-150500.1.1.x86_64.rpm
```

```
[root@worker-01 ~]# sudo yum install -y kubelet kubeadm kubectl --disablereleases=kubernetes
Kubernetes
Last metadata expiration check: 0:00:01 ago on Tue Dec 16 16:11:07 2025.
Dependencies resolved.
=====
 Package          Architecture      Version
=====
Installing:
 kubeadm          x86_64          1.34.3-150500.1.1
 kubectl          x86_64          1.34.3-150500.1.1
 kubelet          x86_64          1.34.3-150500.1.1
Installing dependencies:
 cri-tools        x86_64          1.34.0-150500.1.1
 kubernetes-cni   x86_64          1.7.1-150500.1.1
Transaction summary
=====
Install 5 Packages

Total download size: 52 M
Installed size: 282 M
Downloading Packages:
(1/5): cri-tools-1.34.0-150500.1.1.x86_64.rpm
(2/5): kubeadm-1.34.3-150500.1.1.x86_64.rpm
(3/5): kubectl-1.34.3-150500.1.1.x86_64.rpm
(4/5): kubelet-1.34.3-150500.1.1.x86_64.rpm
(5/5): kubernetes-cni-1.7.1-150500.1.1.x86_64.rpm
-----
Total
Kubernetes
Importing GPG key 0xA296436:
 Userid      : "isv:kubernetes OBS Project <isv:kubernetes@build.opensuse.org>"
 Fingerprint: DE15 B144 86CD 377B 9E87 6E1A 2346 54DA 9A29 6436
 From       : https://pkgs.k8s.io/core:/stable:/v1.34/rpm/repodata/repomd.xml.key
```

```
[root@worker-02 ~]# sudo yum install -y kubelet kubeadm kubectl --disablereleases=kubernetes
Kubernetes
Last metadata expiration check: 0:00:01 ago on Tue Dec 16 16:11:17 2025.
Dependencies resolved.
=====
 Package          Architecture      Version
=====
Installing:
 kubeadm          x86_64          1.34.3-150500.1.1
 kubectl          x86_64          1.34.3-150500.1.1
 kubelet          x86_64          1.34.3-150500.1.1
Installing dependencies:
 cri-tools        x86_64          1.34.0-150500.1.1
 kubernetes-cni   x86_64          1.7.1-150500.1.1
Transaction summary
=====
Install 5 Packages

Total download size: 52 M
Installed size: 282 M
Downloading Packages:
(1/5): kubeadm-1.34.3-150500.1.1.x86_64.rpm
(2/5): cri-tools-1.34.0-150500.1.1.x86_64.rpm
(3/5): kubectl-1.34.3-150500.1.1.x86_64.rpm
(4/5): kubernetes-cni-1.7.1-150500.1.1.x86_64.rpm
(5/5): kubelet-1.34.3-150500.1.1.x86_64.rpm
-----
Total
Kubernetes
Importing GPG key 0xA296436:
 Userid      : "isv:kubernetes OBS Project <isv:kubernetes@build.opensuse.org>"
```

## Enable kubeadm in 3 servers.

- **sudo systemctl enable --now kubelet**

```
[root@master ~]# sudo systemctl enable --now kubelet
Created symlink /etc/systemd/system/multi-user.target.wants/kubelet.service → /usr/lib/systemd/system/kubelet.service.
[root@master ~]# |
```

```
[root@worker-01 ~]# sudo systemctl enable --now kubelet
Created symlink /etc/systemd/system/multi-user.target.wants/kubelet.service → /usr/lib/systemd/system/kubelet.service.
[root@worker-01 ~]#
```

```
[root@worker-02 ~]# sudo systemctl enable --now kubelet
Created symlink /etc/systemd/system/multi-user.target.wants/kubelet.service → /usr/lib/systemd/system/kubelet.service.
[root@worker-02 ~]#
```

Give pod range as like this command only in master machine.

- **kubeadm init --pod-network-cidr=10.244.0.0/16**

```
[root@master ~]# kubeadm init --pod-network-cidr=10.244.0.0/16
[init] Using Kubernetes version: v1.34.3
[preflight] Running pre-flight checks
    [WARNING Hostname]: hostname "master" could not be reached
    [WARNING Hostname]: hostname "master": lookup master on 172.31.0.2:53:
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your in
[preflight] You can also perform this action beforehand using 'kubeadm config
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
[certs] apiserver serving cert is signed for DNS names [kubernetes kubernetes.
ster.local master] and IPs [10.96.0.1 172.31.24.182]
[certs] Generating "apiserver-kubelet-client" certificate and key
[certs] Generating "front-proxy-ca" certificate and key
[certs] Generating "front-proxy-client" certificate and key
[certs] Generating "etcd/ca" certificate and key
[certs] Generating "etcd/server" certificate and key
[certs] etcd/server serving cert is signed for DNS names [localhost master] an
[certs] Generating "etcd/peer" certificate and key
[certs] etcd/peer serving cert is signed for DNS names [localhost master] and
[certs] Generating "etcd/healthcheck-client" certificate and key
[certs] Generating "apiserver-etcd-client" certificate and key
[certs] Generating "sa" key and public key
[kubeconfig] Using kubeconfig folder "/etc/kubernetes"
[kubeconfig] Writing "admin.conf" kubeconfig file
[kubeconfig] Writing "super-admin.conf" kubeconfig file
[kubeconfig] Writing "kubelet.conf" kubeconfig file
```

```
[bootstrap-token] Configured RBAC rules to allow the csrapprover controller automatically approve CSRs from a Node
[bootstrap-token] Configured RBAC rules to allow certificate rotation for all node client certificates in the clus
[bootstrap-token] Creating the "cluster-info" configMap in the "kube-public" namespace
[kubelet-finalize] updating "/etc/kubernetes/kubelet.conf" to point to a rotatable kubelet client certificate and
[addons] Applied essential addon: coreDNS
[addons] Applied essential addon: kube-proxy
```

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Alternatively, if you are the root user, you can run:

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:  
<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 172.31.24.182:6443 --token zylouh.pvmh51kesury66hr \
    --discovery-token-ca-cert-hash sha256:f5c8ca6bbb7f37d1275ed68c3f8fab64700287379194a0655b6ce4dc29e70d19
[root@master ~]#
```

Execute this command in master machine.

- **mkdir -p \$HOME/.kube**
- **sudo cp -i /etc/kubernetes/admin.conf \$HOME/.kube/config**
- **sudo chown \$(id -u):\$(id -g) \$HOME/.kube/config**

```
[root@master ~]# mkdir -p $HOME/.kube
  sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
  sudo chown $(id -u):$(id -g) $HOME/.kube/config
[root@master ~]#
```

Execute this command in 3 servers to join into the cluster.

- **kubeadm join 172.31.24.182:6443 --token zy1ouh.pvmh5lkesury66hr \ --discovery-token-ca-cert-hash sha256:f5c8ca6bbb7f37d1275ed68c3f8fab64700287379194a0655b6ce4dc29e70d19**  
**9194a0655b6ce4dc29e70d19**

```
[root@master ~]# kubeadm join 172.31.24.182:6443 --token zy1ouh.pvmh5lkesury66hr \
  --discovery-token-ca-cert-hash sha256:f5c8ca6bbb7f37d1275ed68c3f8fab64700287379194a0655b6ce4dc29e70d19
[preflight] Running pre-flight checks
  [WARNING Hostname]: hostname "master" could not be reached
  [WARNING Hostname]: hostname "master": lookup master on 172.31.0.2:53: no such host
[preflight] Some fatal errors occurred:
  [ERROR FileAvailable--etc-kubernetes-kubelet.conf]: /etc/kubernetes/kubelet.conf already exists
  [ERROR Port-10250]: Port 10250 is in use
  [ERROR FileAvailable--etc-kubernetes-pki-ca.crt]: /etc/kubernetes/pki/ca.crt already exists
[preflight] If you know what you are doing, you can make a check non-fatal with `--ignore-preflight-errors=...` error: error execution phase preflight: preflight checks failed
To see the stack trace of this error execute with --v=5 or higher
```

```
[root@worker-01 ~]# kubeadm join 172.31.24.182:6443 --token zy1ouh.pvmh5lkesury66hr \
  --discovery-token-ca-cert-hash sha256:f5c8ca6bbb7f37d1275ed68c3f8fab64700287379194a0655b6ce4dc29e70d19
[preflight] Running pre-flight checks
  [WARNING Hostname]: hostname "worker-01" could not be reached
  [WARNING Hostname]: hostname "worker-01": lookup worker-01 on 172.31.0.2:53: no such host
[preflight] Reading configuration from the "kubeadm-config" configMap in namespace "kube-system"...
[preflight] use 'kubeadm init phase upload-config kubeadm --config your-config-file' to re-upload it.
[kubelet-start] writing kubelet configuration to file "/var/lib/kubelet/instance-config.yaml"
[patches] Applied patch of type "application/strategic-merge-patch+json" to target "kubeletconfiguration"
[kubelet-start] writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 501.684629ms
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
```

```
[root@worker-02 ~]# kubeadm join 172.31.24.182:6443 --token zylouh.pvmh51kesury66hr \
--discovery-token-ca-cert-hash sha256:f5c8ca6bbb7f37d1275ed68c3f8fab64700287379194a0655b6ce4dc29e70d19
[preflight] Running pre-flight checks
  [WARNING Hostname]: hostname "worker-02" could not be reached
  [WARNING Hostname]: hostname "worker-02": lookup worker-02 on 172.31.0.2:53: no such host
[preflight] Reading configuration from the "kubeadm-config" ConfigMap in namespace "kube-system"...
[preflight] Use 'kubeadm init phase upload-config kubeadm --config your-config-file' to re-upload it.
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/instance-config.yaml"
[patches] Applied patch of type "application/strategic-merge-patch+json" to target "kubeletconfiguration"
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 502.34032ms
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

[root@worker-02 ~]# |
```

Go to master node and see the nodes.

- **kubectl get nodes**

```
[root@master ~]# kubectl get nodes
NAME      STATUS    ROLES      AGE      VERSION
master    NotReady control-plane   15m     v1.34.3
worker-01  NotReady <none>       2m34s   v1.34.3
worker-02  NotReady <none>       2m27s   v1.34.3
[root@master ~]# |
```

To make status ready execute this command in master machine.

- **kubectl apply -f https://github.com/flannel-io/flannel/releases/latest/download/kube-flannel.yml**

```
[root@master ~]# kubectl apply -f https://github.com/flannel-io/flannel/releases/latest/download/kube-flannel.yml
namespace/kube-flannel created
serviceaccount/flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
[root@master ~]# |
```

- **kubectl get nodes**

```
[root@master ~]# kubectl get nodes
NAME        STATUS   ROLES      AGE        VERSION
master      Ready    control-plane   18m      v1.34.3
worker-01   Ready    <none>     5m40s    v1.34.3
worker-02   Ready    <none>     5m33s    v1.34.3
[root@master ~]#
```

### 3. Run one nginx pod.

Excute this command in master machine.

- **kubectl run firstpod --image=nginx**

```
[root@master ~]# kubectl run firstpod --image=nginx
pod/firstpod created
[root@master ~]# kubectl get pods
NAME        READY   STATUS    RESTARTS   AGE
firstpod   1/1     Running   0          35s
[root@master ~]#
```

- **kubectl expose pod firstpod --port=80 --type=NodePort**
- **kubectl get svc firstpod**

```
[root@master ~]# kubectl get pods
NAME        READY   STATUS    RESTARTS   AGE
firstpod   1/1     Running   0          35s
[root@master ~]# kubectl expose pod firstpod --port=80 --type=NodePort
service/firstpod exposed
[root@master ~]# kubectl get svc firstpod
NAME        TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)        AGE
firstpod   NodePort    10.96.158.54  <none>        80:31716/TCP   20s
[root@master ~]#
```

copy master public ip and port number 31716

← → ⌛ Not secure 3.81.249.223:31716

Gmail YouTube Maps

## Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to [nginx.org](http://nginx.org).  
Commercial support is available at [nginx.com](http://nginx.com).

*Thank you for using nginx.*

## 4. Mug up Master and slave components on k8s.

### MASTER (Control Plane) – “Brain”

#### 4 MAIN COMPONENTS

##### 1. API Server

- Entry point to Kubernetes
- All commands go through this

Think: Gate / Front door

##### 2. etcd

- Stores all cluster data
- Pods, nodes, configs, secrets

Think: Database / Memory

##### 3. Scheduler

- Decides which node a pod runs on

Think: Pod → Node allocator

##### 4. Controller Manager

#### Kubernetes Task - 01

- Keeps desired state = actual state
- Creates/deletes pods if needed

Think: Auto-corrector

## **MASTER IN ONE-LINE MUG UP**

API - ETCD - Scheduler - Controller

SLAVE / WORKER NODE – “Doer”

## **3 MAIN COMPONENTS**

### **1. kubelet**

- Runs on every node
- Talks to master
- Starts pods

Think: Node manager

### **2. Container Runtime**

- Runs containers
- Example: containerd, CRI-O

Think: Engine

### **3. kube-proxy**

- Handles networking
- Service → Pod traffic

Think: Traffic police

## **WORKER IN ONE-LINE MUG UP**

Kubelet – Runtime – Kube-proxy