

1. Deploy an application using a Deployment with 3 replicas and a rolling update strategy.

Create a yaml file.

- **vi rollingupdate.yml**

apiVersion: apps/v1

kind: Deployment

metadata:

name: my-app-deployment

spec:

replicas: 3

strategy:

type: RollingUpdate

rollingUpdate:

maxSurge: 1

maxUnavailable: 1

selector:

matchLabels:

app: my-app

template:

metadata:

labels:

app: my-app

spec:

containers:

- name: my-app-container

image: nginx:latest

ports:

- containerPort: 80

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-app-deployment
spec:
  replicas: 3
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 1
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
  spec:
    containers:
      - name: my-app-container
        image: nginx:latest
        ports:
          - containerPort: 80
```

- **kubectl apply -f rollingupdate.yml**
- **kubectl rollout status deployment my-app-deployment**
- **kubectl get pods -l app=my-app**

```
[root@master ~]# vi rollingupdate.yml
[root@master ~]# kubectl apply -f rollingupdate.yml
deployment.apps/my-app-deployment created
[root@master ~]# kubectl rollout status deployment my-app-deployment
deployment "my-app-deployment" successfully rolled out
[root@master ~]# kubectl get deployment my-app-deployment
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
my-app-deployment   3/3     3            3           110s
[root@master ~]# kubectl get pods -l app=my-app
NAME                           READY   STATUS    RESTARTS   AGE
my-app-deployment-9b8bb4c4-j7rhc   1/1     Running   0          2m9s
my-app-deployment-9b8bb4c4-xkt67   1/1     Running   0          2m9s
my-app-deployment-9b8bb4c4-xwv26   1/1     Running   0          2m9s
[root@master ~]# |
```

2. Configure a Deployment with a Recreate strategy and observe the downtime.

- **vi nginx-recreate.yaml**

apiVersion: apps/v1

kind: Deployment

metadata:

name: nginx-recreate

spec:

replicas: 3

strategy:

type: Recreate

selector:

matchLabels:

app: nginx

template:

metadata:

labels:

app: nginx

spec:

containers:

- name: nginx

image: nginx:1.24

ports:

- containerPort: 80

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-recreate
spec:
  replicas: 3
  strategy:
    type: Recreate
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.24
          ports:
            - containerPort: 80
```

- **kubectl apply -f nginx-recreate.yaml**
- **kubectl get deploy**
- **kubectl get pods**

```
[root@master ~]# vi nginx-recreate.yaml
[root@master ~]# kubectl apply -f nginx-recreate.yaml
deployment.apps/nginx-recreate created
[root@master ~]# kubectl get deploy
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
nginx-recreate 3/3     3            3           12s
[root@master ~]# kubectl get pods
NAME                           READY   STATUS    RESTARTS   AGE
nginx-recreate-6d8cd7bf4b-418vv 1/1     Running   0          22s
nginx-recreate-6d8cd7bf4b-gm6x5 1/1     Running   0          22s
nginx-recreate-6d8cd7bf4b-ppv8z 1/1     Running   0          22s
[root@master ~]#
```

- **kubectl expose deployment nginx-recreate --type=NodePort --port=80**
- **kubectl get svc nginx-recreate**

```
[root@master ~]# kubectl expose deployment nginx-recreate --type=NodePort --port=80
service/nginx-recreate exposed
[root@master ~]# kubectl get svc nginx-recreate
NAME      TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
nginx-recreate   NodePort    10.96.179.144   <none>        80:30277/TCP   10s
[root@master ~]#
```

if you check with ip addresses and with portnumber 30277



Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.



Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.

Keep as watching the pods running by using

- **kubectl get pods -w**

```
[root@master ~]# kubectl get pods -w
NAME                      READY   STATUS    RESTARTS   AGE
nginx-recreate-6d8cd7bf4b-418vv   1/1     Running   0          6m11s
nginx-recreate-6d8cd7bf4b-gm6x5   1/1     Running   0          6m11s
nginx-recreate-6d8cd7bf4b-ppv8z   1/1     Running   0          6m11s
```

open another terminal and login same master machine and perform this.

- **kubectl set image deployment/nginx-recreate nginx=nginx:1.25**

```
[root@master ~]# kubectl get pods -w
NAME          READY   STATUS    RESTARTS   AGE
nginx-recreate-6d8cd7bf4b-418vv  1/1    Running   0          6m11s
nginx-recreate-6d8cd7bf4b-gm6x5  1/1    Running   0          6m11s
nginx-recreate-6d8cd7bf4b-ppv8z  1/1    Running   0          6m11s
nginx-recreate-6d8cd7bf4b-gm6x5  1/1    Terminating   0          6m11s
nginx-recreate-6d8cd7bf4b-418vv  1/1    Terminating   0          6m11s
nginx-recreate-6d8cd7bf4b-ppv8z  1/1    Terminating   0          6m11s
nginx-recreate-6d8cd7bf4b-gm6x5  1/1    Terminating   0          6m11s
nginx-recreate-6d8cd7bf4b-ppv8z  1/1    Terminating   0          6m11s
nginx-recreate-6d8cd7bf4b-gm6x5  1/1    Terminating   0          6m11s
nginx-recreate-6d8cd7bf4b-ppv8z  1/1    Terminating   0          6m11s
nginx-recreate-6d8cd7bf4b-418vv  1/1    Terminating   0          6m11s
nginx-recreate-6d8cd7bf4b-ppv8z  1/1    Terminating   0          6m11s
nginx-recreate-6d8cd7bf4b-gm6x5  1/1    Completed   0          6m11s
nginx-recreate-6d8cd7bf4b-ppv8z  1/1    Completed   0          6m11s
nginx-recreate-6d8cd7bf4b-gm6x5  1/1    Completed   0          6m11s
nginx-recreate-6d8cd7bf4b-418vv  1/1    Completed   0          6m11s
nginx-recreate-77bf8679f9-cfkxp  0/1    Pending    0          6m11s
nginx-recreate-77bf8679f9-cfkxp  0/1    Pending    0          6m11s
nginx-recreate-77bf8679f9-kbkg5  0/1    Pending    0          6m11s
nginx-recreate-77bf8679f9-hf6cw  0/1    Pending    0          6m11s
nginx-recreate-77bf8679f9-kbkg5  0/1    Pending    0          6m11s
nginx-recreate-77bf8679f9-hf6cw  0/1    Pending    0          6m11s
nginx-recreate-77bf8679f9-cfkxp  0/1    ContainerCreating  0          6m11s
nginx-recreate-77bf8679f9-kbkg5  0/1    ContainerCreating  0          6m11s
nginx-recreate-77bf8679f9-hf6cw  0/1    ContainerCreating  0          6m11s
nginx-recreate-6d8cd7bf4b-gm6x5  0/1    Completed   0          6m11s
nginx-recreate-6d8cd7bf4b-ppv8z  0/1    Completed   0          6m11s
nginx-recreate-6d8cd7bf4b-418vv  0/1    Completed   0          6m11s
nginx-recreate-6d8cd7bf4b-418vv  0/1    Completed   0          6m11s
nginx-recreate-77bf8679f9-hf6cw  0/1    Running    0          6m11s
nginx-recreate-77bf8679f9-cfkxp  0/1    Running    0          6m11s
nginx-recreate-77bf8679f9-kbkg5  0/1    Running    0          6m11s
```

```
[root@master ~]# kubectl get pods -w
NAME          READY   STATUS    RESTARTS   AGE
nginx-recreate-6d8cd7bf4b-418vv  1/1    Running   0          6m11s
nginx-recreate-6d8cd7bf4b-gm6x5  1/1    Running   0          6m11s
nginx-recreate-6d8cd7bf4b-ppv8z  1/1    Running   0          6m11s
nginx-recreate-6d8cd7bf4b-gm6x5  1/1    Terminating   0          10m
nginx-recreate-6d8cd7bf4b-418vv  1/1    Terminating   0          10m
nginx-recreate-6d8cd7bf4b-ppv8z  1/1    Terminating   0          10m
nginx-recreate-6d8cd7bf4b-gm6x5  1/1    Terminating   0          10m
nginx-recreate-6d8cd7bf4b-ppv8z  1/1    Terminating   0          10m
nginx-recreate-6d8cd7bf4b-418vv  1/1    Terminating   0          10m
nginx-recreate-6d8cd7bf4b-ppv8z  0/1    Completed   0          10m
nginx-recreate-6d8cd7bf4b-gm6x5  0/1    Completed   0          10m
nginx-recreate-6d8cd7bf4b-418vv  0/1    Completed   0          10m
nginx-recreate-77bf8679f9-cfkxp  0/1    Pending    0          10m
nginx-recreate-77bf8679f9-cfkxp  0/1    Pending    0          10m
nginx-recreate-77bf8679f9-kbkg5  0/1    Pending    0          10m
nginx-recreate-77bf8679f9-hf6cw  0/1    Pending    0          10m
nginx-recreate-77bf8679f9-kbkg5  0/1    Pending    0          10m
nginx-recreate-77bf8679f9-hf6cw  0/1    Pending    0          10m
nginx-recreate-77bf8679f9-cfkxp  0/1    ContainerCreating  0          10m
nginx-recreate-77bf8679f9-kbkg5  0/1    ContainerCreating  0          10m
nginx-recreate-77bf8679f9-hf6cw  0/1    ContainerCreating  0          10m
nginx-recreate-6d8cd7bf4b-gm6x5  0/1    Completed   0          10m
nginx-recreate-6d8cd7bf4b-ppv8z  0/1    Completed   0          10m
nginx-recreate-6d8cd7bf4b-ppv8z  0/1    Completed   0          10m
nginx-recreate-6d8cd7bf4b-418vv  0/1    Completed   0          10m
nginx-recreate-6d8cd7bf4b-418vv  0/1    Completed   0          10m
nginx-recreate-77bf8679f9-hf6cw  1/1    Running    0          10m
nginx-recreate-77bf8679f9-cfkxp  1/1    Running    0          10m
nginx-recreate-77bf8679f9-kbkg5  1/1    Running    0          10m
^C[root@master ~]# kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
nginx-recreate-77bf8679f9-cfkxp  1/1    Running   0          82s
nginx-recreate-77bf8679f9-hf6cw  1/1    Running   0          82s
nginx-recreate-77bf8679f9-kbkg5  1/1    Running   0          82s
[root@master ~]
```



3. Update an existing Deployment and perform a rollback to the previous version.

- **kubectl get deployments**
- **kubectl get deployments nginx-recreate -o wide**

```
[root@master ~]# kubectl get deployments
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
nginx-recreate   3/3     3           3          20m
[root@master ~]# kubectl get deployments nginx-recreate -o wide
NAME        READY   UP-TO-DATE   AVAILABLE   AGE   CONTAINERS   IMAGES      SELECTOR
nginx-recreate   3/3     3           3          20m   nginx       nginx:1.25   app=nginx
```

- **kubectl rollout history deployment nginx-recreate**

```
[root@master ~]# kubectl rollout history deployment nginx-recreate
deployment.apps/nginx-recreate
REVISION  CHANGE-CAUSE
1          <none>
2          <none>
```

- **kubectl rollout undo deployment nginx-recreate**

```
[root@master ~]# kubectl rollout undo deployment nginx-recreate
deployment.apps/nginx-recreate rolled back
[root@master ~]# kubectl get pods -w
NAME                               READY   STATUS    RESTARTS   AGE
nginx-recreate-6d8cd7bf4b-7jlsr   1/1    Running   0          29s
nginx-recreate-6d8cd7bf4b-m5h87   1/1    Running   0          29s
nginx-recreate-6d8cd7bf4b-vdhvd   1/1    Running   0          29s
```

- **kubectl get deployment nginx-recreate -o wide**

```
[root@master ~]# kubectl get deployment nginx-recreate -o wide
NAME           READY   UP-TO-DATE   AVAILABLE   AGE   CONTAINERS   IMAGES     SELEC
TOR
nginx-recreate 3/3      3           3           22m   nginx        nginx:1.24   app=n
ginx
[root@master ~]# kubectl get pods
NAME                               READY   STATUS    RESTARTS   AGE
nginx-recreate-6d8cd7bf4b-7jlsr   1/1    Running   0          92s
nginx-recreate-6d8cd7bf4b-m5h87   1/1    Running   0          92s
nginx-recreate-6d8cd7bf4b-vdhvd   1/1    Running   0          92s
[root@master ~]# kubectl rollout history deployment nginx-recreate
deployment.apps/nginx-recreate
REVISION  CHANGE-CAUSE
2          <none>
3          <none>
```

4. Modify a Deployment to add resource requests and limits for CPU and memory.

- **vi nginx-resources.yaml**

apiVersion: apps/v1

kind: Deployment

metadata:

name: nginx-resources

spec:

replicas: 3

strategy:

type: RollingUpdate

selector:

matchLabels:

app: nginx

template:

metadata:

labels:

app: nginx

spec:

containers:

- name: nginx

image: nginx:1.25

ports:

- containerPort: 80

resources:

requests:

cpu: "100m"

memory: "128Mi"

limits:

cpu: "200m"

memory: "256Mi"

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-resources
spec:
  replicas: 3
  strategy:
    type: RollingUpdate
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.25
          ports:
            - containerPort: 80
          resources:
            requests:
              cpu: "100m"
              memory: "128Mi"
            limits:
              cpu: "200m"
              memory: "256Mi"
```

- **kubectl apply -f nginx-resources.yaml**
- **kubectl get deployment nginx-resources**
- **kubectl get pods**

```
[root@master ~]# vi nginx-resources.yaml
[root@master ~]# kubectl apply -f nginx-resources.yaml
deployment.apps/nginx-resources created
[root@master ~]# kubectl get deployment nginx-resources
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
nginx-resources   3/3     3            3           17s
[root@master ~]# kubectl get pods
NAME                           READY   STATUS    RESTARTS   AGE
nginx-resources-fb5cb79f5-2zvbr   1/1     Running   0          31s
nginx-resources-fb5cb79f5-rzsxk   1/1     Running   0          31s
nginx-resources-fb5cb79f5-sv48l   1/1     Running   0          31s
```

- **kubectl describe pod nginx-resources-fb5cb79f5-2zvbr**

```
[root@master ~]# kubectl describe pod nginx-resources-fb5cb79f5-2zvbr
Name:           nginx-resources-fb5cb79f5-2zvbr
Namespace:      default
Priority:       0
Service Account: default
Node:          worker-01/172.31.27.50
Start Time:    Sat, 20 Dec 2025 16:30:22 +0000
Labels:         app=nginx
                pod-template-hash=fb5cb79f5
Annotations:   <none>
Status:        Running
IP:            10.244.1.32
IPs:
  IP:          10.244.1.32
Controlled By: ReplicaSet/nginx-resources-fb5cb79f5
Containers:
  nginx:
    Container ID:  containerd://60c0ecce9198c38620b608fe536221dd3952c0285f0eea8e1
    Image:         nginx:1.25
    Image ID:     docker.io/library/nginx@sha256:a484819eb60211f5299034ac80f6a68
    Port:          80/TCP
    Host Port:    0/TCP
    State:        Running
      Started:   Sat, 20 Dec 2025 16:30:23 +0000
    Ready:        True
    Restart Count: 0
    Limits:
      cpu:        200m
      memory:    256Mi
    Requests:
      cpu:        100m
      memory:    128Mi
    Environment: <none>
    Mounts:
```

5. Create a Deployment with MaxSurge and MaxUnavailable configurations.

- **Vi nginx-rolling.yaml**

apiVersion: apps/v1

kind: Deployment

metadata:

name: nginx-rolling

spec:

```
replicas: 4
strategy:
  type: RollingUpdate
  rollingUpdate:
    maxSurge: 1
    maxUnavailable: 1
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
        image: nginx:1.25
      ports:
        - containerPort: 80
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-rolling
spec:
  replicas: 4
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 1
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
  spec:
    containers:
      - name: nginx
        image: nginx:1.25
        ports:
          - containerPort: 80
```

- **kubectl apply -f nginx-rolling.yaml**
- **kubectl get deployment nginx-rolling**
- **kubectl describe deployment nginx-rolling**

```
[root@master ~]# vi nginx-rolling.yaml
[root@master ~]# kubectl apply -f nginx-rolling.yaml
deployment.apps/nginx-rolling created
[root@master ~]# kubectl get deployment nginx-rolling
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
nginx-rolling  4/4     4            4           9s
[root@master ~]# kubectl describe deployment nginx-rolling
Name:           nginx-rolling
Namespace:      default
CreationTimestamp: Sat, 20 Dec 2025 16:40:30 +0000
Labels:         <none>
Annotations:    deployment.kubernetes.io/revision: 1
Selector:       app=nginx
Replicas:       4 desired | 4 updated | 4 total | 4 available | 0 unavailable
StrategyType:   RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 1 max unavailable, 1 max surge
Pod Template:
  Labels:  app=nginx
  Containers:
    nginx:
      Image:      nginx:1.25
      Port:       80/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
      Node-Selectors: <none>
```

- **kubectl get pods -w**

observe the behaviour it will terminate one and create one.

- **kubectl set image deployment/nginx-rolling**

nginx=nginx:1.26

```
root@master:~# kubectl get pods -w
nginx-rolling-77bf8679f9-btdgh 1/1 Running 0 45s
nginx-rolling-77bf8679f9-g6h5z 1/1 Running 0 45s
nginx-rolling-77bf8679f9-k54w4 1/1 Running 0 45s
nginx-rolling-77bf8679f9-kpj55 1/1 Running 0 45s
nginx-rolling-57489d7c8d-2mg6c 0/1 Pending 0 0s
nginx-rolling-57489d7c8d-2mg6c 0/1 Pending 0 0s
nginx-rolling-77bf8679f9-kmg6c 0/1 Pending 0 0s
nginx-rolling-77bf8679f9-k54w4 1/1 Terminating 0 2m
nginx-rolling-57489d7c8d-2mg6c 0/1 ContainerCreating 0
nginx-rolling-57489d7c8d-jggtv 0/1 Pending 0
nginx-rolling-77bf8679f9-k54w4 1/1 Terminating 0
nginx-rolling-57489d7c8d-jggtv 0/1 Pending 0
nginx-rolling-57489d7c8d-jggtv 0/1 ContainerCreating 0
nginx-rolling-77bf8679f9-k54w4 0/1 Completed 0
nginx-rolling-77bf8679f9-k54w4 0/1 Completed 0
nginx-rolling-77bf8679f9-k54w4 0/1 Completed 0
nginx-rolling-57489d7c8d-2mg6c 1/1 Running 0
nginx-rolling-77bf8679f9-kpj55 1/1 Terminating 0
nginx-rolling-57489d7c8d-d656z 0/1 Pending 0
nginx-rolling-77bf8679f9-kpj55 1/1 Terminating 0
nginx-rolling-57489d7c8d-d656z 0/1 Pending 0
nginx-rolling-57489d7c8d-jggtv 0/1 ContainerCreating 0
nginx-rolling-57489d7c8d-jggtv 1/1 Running 0
nginx-rolling-77bf8679f9-g6h5z 1/1 Terminating 0
nginx-rolling-57489d7c8d-wx5qq 0/1 Pending 0
nginx-rolling-77bf8679f9-g6h5z 1/1 Terminating 0
nginx-rolling-57489d7c8d-wx5qq 0/1 Pending 0
nginx-rolling-57489d7c8d-wx5qq 0/1 ContainerCreating 0
nginx-rolling-77bf8679f9-kpj55 0/1 Completed 0
nginx-rolling-57489d7c8d-d656z 1/1 Running 0
nginx-rolling-77bf8679f9-btdgh 1/1 Terminating 0
nginx-rolling-77bf8679f9-btdgh 1/1 Terminating 0
nginx-rolling-77bf8679f9-g6h5z 0/1 Completed 0
nginx-rolling-57489d7c8d-wx5qq 1/1 Running 0
nginx-rolling-77bf8679f9-kpj55 0/1 Completed 0
```

```
root@master:~# kubectl set image deployment/nginx-rolling nginx=nginx:1.26
deployment.apps/nginx-rolling image updated
[root@master ~]#
```

nginx-rolling-77bf8679f9-dtgh	1/1	Running	0	45s
nginx-rolling-77bf8679f9-g6h5z	1/1	Running	0	45s
nginx-rolling-77bf8679f9-k54w4	1/1	Running	0	45s
nginx-rolling-77bf8679f9-kpj55	1/1	Running	0	45s
nginx-rolling-57489d7c8d-2mg6c	0/1	Pending	0	0s
nginx-rolling-57489d7c8d-2mg6c	0/1	Pending	0	0s
nginx-rolling-77bf8679f9-k54w4	1/1	Terminating	0	2m8s
nginx-rolling-57489d7c8d-2mg6c	0/1	ContainerCreating	0	0s
nginx-rolling-57489d7c8d-jggtv	0/1	Pending	0	0s
nginx-rolling-77bf8679f9-k54w4	1/1	Terminating	0	2m9s
nginx-rolling-57489d7c8d-jggtv	0/1	Pending	0	0s
nginx-rolling-57489d7c8d-jggtv	0/1	ContainerCreating	0	0s
nginx-rolling-77bf8679f9-k54w4	0/1	Completed	0	2m9s
nginx-rolling-77bf8679f9-k54w4	0/1	Completed	0	2m9s
nginx-rolling-77bf8679f9-k54w4	0/1	Completed	0	2m9s
nginx-rolling-57489d7c8d-2mg6c	1/1	Running	0	3s
nginx-rolling-77bf8679f9-kpj55	1/1	Terminating	0	2m11s
nginx-rolling-57489d7c8d-d656z	0/1	Pending	0	0s
nginx-rolling-77bf8679f9-kpj55	1/1	Terminating	0	2m11s
nginx-rolling-57489d7c8d-d656z	0/1	Pending	0	0s
nginx-rolling-57489d7c8d-d656z	0/1	ContainerCreating	0	0s
nginx-rolling-57489d7c8d-jggtv	1/1	Running	0	2s
nginx-rolling-77bf8679f9-g6h5z	1/1	Terminating	0	2m11s
nginx-rolling-57489d7c8d-wx5qq	0/1	Pending	0	0s
nginx-rolling-77bf8679f9-g6h5z	1/1	Terminating	0	2m11s
nginx-rolling-57489d7c8d-wx5qq	0/1	Pending	0	0s
nginx-rolling-57489d7c8d-wx5qq	0/1	ContainerCreating	0	1s
nginx-rolling-77bf8679f9-kpj55	0/1	Completed	0	2m12s
nginx-rolling-57489d7c8d-d656z	1/1	Running	0	1s
nginx-rolling-77bf8679f9-btdgh	1/1	Terminating	0	2m12s
nginx-rolling-77bf8679f9-btdgh	1/1	Terminating	0	2m12s
nginx-rolling-77bf8679f9-g6h5z	0/1	Completed	0	2m12s
nginx-rolling-57489d7c8d-wx5qq	1/1	Running	0	1s
nginx-rolling-77bf8679f9-kpj55	0/1	Completed	0	2m12s
nginx-rolling-77bf8679f9-kpj55	0/1	Completed	0	2m12s
nginx-rolling-77bf8679f9-btdgh	0/1	Completed	0	2m12s
nginx-rolling-77bf8679f9-g6h5z	0/1	Completed	0	2m13s

6. Set up a Deployment with a custom revision history limit.

- **vi nginx-revision-limit.yaml**

apiVersion: apps/v1

kind: Deployment

metadata:

name: nginx-revision-limit

spec:

replicas: 3

revisionHistoryLimit: 2

strategy:

type: RollingUpdate

rollingUpdate:

maxSurge: 1

maxUnavailable: 1

selector:

matchLabels:

app: nginx

template:

metadata:

labels:

app: nginx

spec:

containers:

- name: nginx

image: nginx:1.25

ports:

- containerPort: 80

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-revision-limit
spec:
  replicas: 3
  revisionHistoryLimit: 2
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 1
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.25
          ports:
            - containerPort: 80

```

- **kubectl apply -f nginx-revision-limit.yaml**
- **kubectl describe deployment nginx-revision-limit**

```

[root@master ~]# vi nginx-revision-limit.yaml
[root@master ~]# kubectl apply -f nginx-revision-limit.yaml
deployment.apps/nginx-revision-limit created
[root@master ~]# kubectl describe deployment nginx-revision-limit
Name:           nginx-revision-limit
Namespace:      default
CreationTimestamp: Sat, 20 Dec 2025 17:04:16 +0000
Labels:         <none>
Annotations:   deployment.kubernetes.io/revision: 1
Selector:       app=nginx
Replicas:      3 desired | 3 updated | 3 total | 3 available | 0 unavailable
StrategyType:  RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 1 max unavailable, 1 max surge
Pod Template:
  Labels:  app=nginx
  Containers:
    nginx:
      Image:      nginx:1.25
      Port:       80/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
  Node Selector: <none>

```

Update the image multiple times

- **kubectl set image deployment/nginx-revision-limit nginx=nginx:1.26**
- **kubectl set image deployment/nginx-revision-limit nginx=nginx:1.27**
- **kubectl set image deployment/nginx-revision-limit nginx=nginx:1.28**
- **kubectl get rs**

```
[root@master ~]# kubectl set image deployment/nginx-revision-limit nginx=nginx:1.26
deployment.apps/nginx-revision-limit image updated
[root@master ~]# kubectl set image deployment/nginx-revision-limit nginx=nginx:1.27
deployment.apps/nginx-revision-limit image updated
[root@master ~]# kubectl set image deployment/nginx-revision-limit nginx=nginx:1.28
deployment.apps/nginx-revision-limit image updated
[root@master ~]# kubectl get rs
NAME           DESIRED   CURRENT   READY   AGE
nginx-revision-limit-57489d7c8d  0         0         0      36s
nginx-revision-limit-7489fb554f  0         0         0      20s
nginx-revision-limit-7f9c8bf8cd  3         3         3      9s
[root@master ~]# |
```

7. Pause a Deployment during an update and then resume it.

Use nginx-rolling yamlfile and update the image to 1.25

- **kubectl set image deployment/nginx-rolling nginx=nginx:1.25**

```
[root@master ~]# kubectl set image deployment/nginx-rolling nginx=nginx:1.25
deployment.apps/nginx-rolling image updated
```

- **kubectl rollout pause deployment nginx-rolling**

```
[root@master ~]# kubectl rollout pause deployment nginx-rolling
error: deployments.apps "nginx-rolling" is already paused
```

- **kubectl rollout resume deployment nginx-rolling**

```
[root@master ~]# kubectl rollout resume deployment nginx-rolling
deployment.apps/nginx-rolling resumed
```

- **kubectl rollout status deployment nginx-rolling**
- **kubectl get pods -w**

```
[root@master ~]# kubectl rollout status deployment nginx-rolling
deployment "nginx-rolling" successfully rolled out
[root@master ~]# kubectl get pods -w
NAME           READY   STATUS    RESTARTS   AGE
nginx-rolling-77bf8679f9-h6dfk  1/1     Running   0          23s
nginx-rolling-77bf8679f9-hkxnc  1/1     Running   0          22s
nginx-rolling-77bf8679f9-mjxjq  1/1     Running   0          21s
nginx-rolling-77bf8679f9-vcmsf  1/1     Running   0          23s
^C[root@master ~]# kubectl get deployment nginx-rolling -o yaml
NAME      READY   UP-TO-DATE   AVAILABLE   AGE   CONTAINERS   IMAGES   SELECTOR
nginx-rolling  4/4     4           4          9m43s   nginx        nginx:1.25  app=nginx
[root@master ~]#
```

8. Create a pod using resource requests for memory and CPU and observe how the scheduler assigns it to a node.

- **vi resource-pod.yaml**

apiVersion: v1

kind: Pod

metadata:

name: resource-pod

spec:

containers:

- name: nginx

image: nginx

resources:

requests:

cpu: "500m"

memory: "256Mi"

limits:

cpu: "1"

memory: "512Mi"

```
apiVersion: v1
kind: Pod
metadata:
  name: resource-pod
spec:
  containers:
  - name: nginx
    image: nginx
    resources:
      requests:
        cpu: "500m"
        memory: "256Mi"
      limits:
        cpu: "1"
        memory: "512Mi"
```

- **kubectl apply -f resource-pod.yaml**
- **kubectl get pod resource-pod -o wide**

```
[root@master ~]# vi resource-pod.yaml
[root@master ~]# kubectl apply -f resource-pod.yaml
pod/resource-pod created
[root@master ~]# kubectl get pod resource-pod -o wide
NAME      READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
resource-pod  1/1    Running   0          12s    10.244.2.51  worker-02  <none>        <none>
```

- **kubectl describe pod resource-pod**

```
[root@master ~]# kubectl describe pod resource-pod
Name:           resource-pod
Namespace:      default
Priority:       0
Service Account: default
Node:          worker-02/172.31.72.26
Start Time:    Sat, 20 Dec 2025 17:36:51 +0000
Labels:         <none>
Annotations:   <none>
Status:        Running
IP:            10.244.2.51
IPs:
  IP:  10.244.2.51
Containers:
  nginx:
    Container ID:  containerd://e208dc5a0a1e325942f8ec25bcc5a0902a155f3a8df58ed11a1
    Image:          nginx
    Image ID:      docker.io/library/nginx@sha256:fb01117203ff38c2f9af91db1a7409459
    Port:          <none>
    Host Port:    <none>
    State:        Running
      Started:    Sat, 20 Dec 2025 17:36:52 +0000
    Ready:        True
    Restart Count: 0
    Limits:
      cpu:        1
      memory:    512Mi
    Requests:
      cpu:        500m
      memory:    256Mi
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-g855m (ro)
Conditions:
  Type          Status
  PodReadyToStartContainers  True
  Initialized   True
```

It was assigned on worker-02

- **kubectl describe node worker-02**

```
[root@master ~]# kubectl describe node worker-02
Name:           worker-02
Roles:          <none>
Labels:         beta.kubernetes.io/arch=amd64
                beta.kubernetes.io/os=linux
                kubernetes.io/arch=amd64
                kubernetes.io/hostname=worker-02
                kubernetes.io/os=linux
Annotations:   flannel.alpha.coreos.com/backend-data: {"VNI":1,"VtepMAC":"02:29:ad:52:c0:48"}
                flannel.alpha.coreos.com/backend-type: vxlan
                flannel.alpha.coreos.com/kube-subnet-manager: true
                flannel.alpha.coreos.com/public-ip: 172.31.72.26
                node.alpha.kubernetes.io/ttl: 0
                volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp: Tue, 16 Dec 2025 16:32:11 +0000
Taints:         <none>
Unschedulable:  false
Lease:
  HolderIdentity:  worker-02
  AcquireTime:    <unset>
  RenewTime:     Sat, 20 Dec 2025 17:38:03 +0000
Conditions:
  Type          Status  LastHeartbeatTime          LastTransitionTime          Reason
  ----          ----   -----                      -----                  -----
  NetworkUnavailable  False   Sat, 20 Dec 2025 15:35:03 +0000  Sat, 20 Dec 2025 15:35:03 +0000  Flanne
g on this node
  MemoryPressure   False   Sat, 20 Dec 2025 17:36:37 +0000  Fri, 19 Dec 2025 15:00:46 +0000  Kubelet
client memory available
  DiskPressure    False   Sat, 20 Dec 2025 17:36:37 +0000  Fri, 19 Dec 2025 15:00:46 +0000  Kubelet
sk pressure
  PIDPressure    False   Sat, 20 Dec 2025 17:36:37 +0000  Fri, 19 Dec 2025 15:00:46 +0000  Kubelet
client PID available
  Ready          True    Sat, 20 Dec 2025 17:36:37 +0000  Fri, 19 Dec 2025 15:00:46 +0000  Kubelet
```

```

pods:           110
Allocatable:
  cpu:          2
  ephemeral-storage: 19255611770
  hugepages-1Gi:   0
  hugepages-2Mi:   0
  memory:        3808856Ki
  pods:          110
System Info:
  Machine ID:      ec20e49319d0cb0218d9174bf8e17800
  System UUID:     ec20e493-19d0-cb02-18d9-174bf8e17800
  Boot ID:        7a470162-2749-45b3-823e-a73942549657
  Kernel Version:  6.1.158-180.294.amzn2023.x86_64
  OS Image:       Amazon Linux 2023.9.20251208
  Operating System: linux
  Architecture:    amd64
  Container Runtime Version: containerd://2.1.5
  Kubelet Version: v1.34.3
  Kube-Proxy Version:
PodCIDR:          10.244.2.0/24
PodCIDRs:         10.244.2.0/24
Non-terminated Pods: (3 in total)
  Namespace          Name            CPU Requests  CPU Limits  Memory Requests  Memory Limits  Age
  ----              ---            -----          -----        -----          -----          ---
  default           resource-pod      500m (25%)  1 (50%)    256Mi (6%)    512Mi (13%)  81s
  kube-flannel      kube-flannel-ds-chq7f  100m (5%)   0 (0%)    50Mi (1%)    0 (0%)    4d1h
  kube-system       kube-proxy-9v7j7      0 (0%)      0 (0%)    0 (0%)      0 (0%)    4d1h
Allocated resources:
(Total limits may be over 100 percent, i.e., overcommitted.)
Resource      Requests  Limits
-----          -----  -----
cpu            600m (30%) 1 (50%)
memory        306Mi (8%) 512Mi (13%)
ephemeral-storage 0 (0%) 0 (0%)
hugepages-1Gi  0 (0%) 0 (0%)
hugepages-2Mi  0 (0%) 0 (0%)
Events: <none>
[root@master ~]# 

```